

Jemin Woo

Professional Experience | 2026

AWS GPUaaS Project

May 2026 - Present

Achievements

- Provided technical guidance and performed hands-on cloud engineering work to migrate GPU workloads for AI training to AWS infrastructure.
- Served as project manager within the system engineering organization, successfully managing resource allocation and project schedules.

Technical skills: AWS SageMaker HyperPod | AWS SageMaker HyperPod with Amazon EKS | Slurm | GPU Training

Server Provisioning and Operations Platform Development

January 2026 - Present

Achievements

- Planned and developed a service for server provisioning requests, operations, and management.
- Led the project, defined work items, and assigned team members based on technical expertise, successfully launching the service on schedule.

Technical skills: Next.js

NVIDIA RTX PRO 6000 Evaluation PoC

January 2026 - Present

Achievements

- Conducted infrastructure compatibility and performance testing to evaluate NVIDIA RTX PRO 6000 ahead of adoption.
- Ran MLPerf Inference benchmarks to validate inference workload performance.

Technical skills: MLPerf Inference | NVIDIA GPU

Professional Experience

NVIDIA B200 InfiniBand Cluster Deployment Support

September 2025 - January 2026

Achievements

- Participated in a government GPU initiative and provided technical support for the deployment of an NVIDIA B200 InfiniBand cluster.

Technical skills: InfiniBand | NVIDIA GPU | Slurm | GPU Training

NVIDIA B200 Evaluation PoC

September 2025 - November 2025

Achievements

- Conducted infrastructure compatibility and performance testing to evaluate NVIDIA B200 ahead of adoption.
- Ran MLPerf Inference and MLPerf Training benchmarks to validate inference and training workload performance.

Technical skills: MLPerf Inference | MLPerf Training | NVIDIA GPU | GPU Training

NVIDIA H200 RoCEv2 Cluster Deployment Support

December 2024 - April 2025

Achievements

- Successfully deployed a GPU cluster using NVIDIA H200 systems and RoCEv2.
- Served as technical project manager, directly resolving issues during cluster deployment, managing the schedule, and coordinating with related teams.

Technical skills: NVIDIA GPU | RoCEv2 | Slurm | GPU Training

Professional Experience

Intel Gaudi 3 and AMD MI300X Evaluation PoC

November 2024 - May 2025

Achievements

- Conducted infrastructure compatibility and performance testing to evaluate Intel Gaudi 3 ahead of adoption.
- Conducted infrastructure compatibility and performance testing to evaluate AMD MI300X ahead of adoption.

Responsibilities

- Led the PoC, coordinated with related teams, defined project tasks, and managed the schedule.
- Performed GPU workload testing as the engineer responsible for infrastructure compatibility and performance validation.

Technical skills: MLPerf Training | GPU Training

RoCEv2 GPU Clustering PoC

September 2024 - November 2024

Achievements

- Conducted a PoC for deploying a GPU cluster using RoCEv2 and successfully completed testing.

Responsibilities

- Led the PoC, coordinated with related teams, defined project tasks, and managed the schedule.
- Researched and tested the technologies required to deploy the GPU cluster.
- Documented PoC deliverables and shared them internally.

Technical skills: MLPerf Training | RoCEv2 | GPU Training | Slurm

Professional Experience

AWS S3 API-Compatible Storage PoC

October 2023 - December 2023

Achievements

- Completed functionality and reliability evaluations of AWS S3 API-compatible storage from three vendors at the minimum cost proposed by the vendors.

Responsibilities

- Led the PoC and facilitated communication among DBAs, system engineers, and vendors.
- Contributed to the selection of an S3-compatible storage vendor for SingleStore DB.
- As the system engineer, defined infrastructure reliability test cases and coordinated with data center teams and DBAs to plan and execute testing.
- Communicated DBA requirements to each vendor and coordinated the timely implementation of needed improvements.

Technical skills: AWS S3 API | S3-Compatible Storage | SingleStore DB | Elasticsearch

Kubernetes-Based Elasticsearch PaaS Development

December 2020 - August 2021

Achievements

- Deployed and operated 200 Elasticsearch clusters using a fully automated deployment system.
- Saved more than two person-months of engineering effort by automating deployment and failure recovery.

Responsibilities

- Planned the automated deployment system and configured its core components.
- Developed and operated the web interface and APIs for the PaaS service.
- Deployed and operated Kubernetes clusters.
- Developed OS templates and operated server infrastructure.
- Developed Ansible playbooks for system configuration.
- Conducted benchmark testing of Kubernetes-based Elasticsearch clusters.

Technical skills: Kubernetes | Linux Kickstart | Elasticsearch | ECK Operator | Python | FastAPI | Frontend Development