

JEMIN WOO

Seoul, Republic of Korea | jemini12@naver.com | linkedin.com/in/jemin-woo

PROFESSIONAL SUMMARY

Infrastructure engineer with 10+ years of experience across production server, cloud, Kubernetes, platform, and AI infrastructure, including 7+ years at Kakao Corp. Recent hands-on work includes NVIDIA B200, H200, and RTX PRO 6000 performance testing; InfiniBand and RoCEv2 GPU clusters; Slurm; MLPerf; AWS GPU; and infrastructure automation.

CORE EXPERTISE

AI/GPU Infrastructure: NVIDIA B200, H200, RTX PRO 6000, Intel Gaudi 3, AMD MI300X, MLPerf Training/Inference

Networking & Orchestration: InfiniBand, RoCEv2, Kubernetes, Slurm, ECK, Ansible

Platform & Operations: Linux, Python, System engineering | Cloud: AWS, GCP

Customer Engagement & Leadership: clear technical communication, trusted stakeholder alignment, cross-functional leadership, high ownership

PROFESSIONAL EXPERIENCE

Kakao Corp | System Engineer | Feb 2019 - Present | Seongnam, Republic of Korea

AI Infrastructure & GPU | 2024 - Present

- Validated compatibility and performance across NVIDIA B200, H200, and RTX PRO 6000 platforms using MLPerf Training and Inference methodologies.
- Served as technical PM for an H200 RoCEv2 GPU cluster, maintaining clear stakeholder communication and driving cross-team resolution across server, networking, and vendors.
- Led a RoCEv2 GPU clustering proof of concept and supported an InfiniBand-based Korean government GPU cluster; used Slurm for workload scheduling and cluster validation.
- Led accelerator evaluations for Intel Gaudi 3 and AMD MI300X, covering compatibility, performance.
- Provided technical consulting and project leadership for migration of on premise GPU training workloads to AWS GPUaaS, owning staffing, schedule, and stakeholder alignment.

Platform Engineering & System Engineering | 2019 - Present

- Planned and developed a server provisioning and operations platform to improve repeatability and operational visibility for infrastructure delivery.
- Built a Kubernetes/ECK operator based Elasticsearch PaaS that automated deployment and operations for 200 clusters, saving more than two person-months of effort per month.
- Established operating policies for more than 500 Elasticsearch clusters and improved consistency across a large production estate.

EARLIER EXPERIENCE

SoftLayer, an IBM Company | Server Build Technician | Feb 2018 - Feb 2019 | Pangyo, Republic of Korea

- Built and provisioned data-center servers, supporting hands-on hardware deployment and server lifecycle operations.

AXA Direct Korea | Systems Operations | May 2016 - Dec 2017 | Seoul, Republic of Korea

- Managed call-system infrastructure across Windows, Unix, Avaya, and database environments; performed operational monitoring and log analysis.

EDUCATION

Sogang University | B.S., Computer Science | 2010 - 2016 | Seoul, Republic of Korea

CERTIFICATIONS

- NVIDIA-Certified Associate: AI Infrastructure and Operations - Issued Oct 2025; expires Oct 2027
- Google Professional Cloud Architect - Issued Mar 2026; expires Mar 2028
- Certified Kubernetes Administrator (CKA), The Linux Foundation - Issued Feb 2020
- Certified Kubernetes Application Developer (CKAD), The Linux Foundation - Issued Apr 2022